



DUBAI CAMPUS

F21BC

BIOLOGICALLY INSPIRED COMPUTATION

COURSEWORK 1A

---

# Implementing Gradient Descent

---

*Author:*  
Heba El-Shimy

*Student Number:*  
H00280277

## Contents


<b>1</b>	<b>Introduction</b>	<b>2</b>
<b>2</b>	<b>General Description</b>	<b>2</b>
<b>3</b>	<b>Choice of model parameters</b>	<b>4</b>
<b>4</b>	<b>Results of the algorithm on the test dataset</b>	<b>7</b>
<b>5</b>	<b>Marking Scheme</b>	<b>11</b>
<b>6</b>	<b>Conclusions</b>	<b>12</b>
<b>7</b>	<b>Appendices</b>	<b>13</b>
7.1	Program Code . . . . .	13
7.2	How to run the code . . . . .	21

# 1 Introduction

This report briefly discusses the concept and implementation of one of the common neural networks learning algorithm, that is, gradient descent. The program accompanying this report represents how gradient descent was used to train a single neuron to identify cat images from non-cat images.

# 2 General Description

An artificial neuron is a concept derived from biology, particularly, the neurons in the human brain. The brain's neurons are the smallest computational units that accept multiple inputs from other neurons, make a calculation over the inputs and then produce an output that's conducted or passed on to another neuron, as one of its multiple inputs.



Human neuron structure versus artificial neuron [2]

The focus in this report will be the learning process of a single artificial neuron. The neuron itself is just a storage unit, that holds a computed value. As for the learning process, it is divided into two main parts, the feedforward part and the backpropagation.

The feedforward part is responsible for collecting all input values, adding different weights to them according to their importance in relation to the output, adding

a bias value and finally summing all these values up and passing them through an activation function which defines a threshold that if exceeded by the calculated sum, then the neuron will fire, i.e. return true/1 as output and false otherwise.

The first step of the feedforward process can be written as follows:

$$z = \sum w_i x_i + b$$

where  $w_i$  is the  $i$ th item in the weights vector,  $x_i$  is the  $i$ th item in the input vector (for a single sample) and  $b$  is the bias vector.

The activation function used in this coursework is the *Sigmoid* function, which outputs a value of  $[0,1]$ . It can be written as follows:

$$a(z) = \frac{1}{1 + e^{-z}}$$

After finishing the feedforward step, we will calculate the loss, i.e., how far away was the algorithm from the true value. This is crucial for the model to adjust and find out better values for the weights and bias that enhance the algorithm's accuracy.

The loss function used in this coursework is the *Cross-Entropy* loss. It suits the problem we have in classifying images to cat and non-cat as it is a *binary classification* problem.

$$L(y, a) = -(y \log(a) + (1 - y) \log(1 - a))$$

where  $a$  is the predicted value and  $y$  is the actual value. As we are dealing with  $m$  number of samples, we will need to take the sum of all losses across all samples.

$$J(w, b) = -\frac{1}{m} \sum_{i=1}^m L(y^i, a^i)$$

The second step in the training process is the backpropagation, and it's where the actual "learning" happens. The intuition behind this step is that calculated cost is sent back to the neuron to readjust its weights. The way this is handled is as an optimization problem, where the algorithm will try to minimize the cost and find the optimal weights and bias that gives the best performance.

The algorithm used for this learning process is *Gradient Descent*, which finds the derivative of the cost function and keeps taking steps iteratively towards the a minimum. *Gradient Descent* is a good algorithm to use with *Cross-Entropy* loss function as it is a perfect convex function with only one minimum, and since *Gradient Descent* is a local search algorithm, we can be sure that it will find the right solution and not be stuck into a local minimum while there is possible other global minimum as can be seen in other loss functions.

### 3 Choice of model parameters

The selection of the algorithm's parameters have been done after multiple experiments with different sets of parameters. Below is a table listing the different experiments that have been done along with the corresponding parameters and the final result represented as the accuracy.

#### - Weights initialization:

I tried different methods for initializing the weights vector by creating random values for the weights. Some of methods I followed were suggested by research and some were random trials by combining those suggestions. The suggested methods by research were:

- a) Random values between -0.5 and 0.5[4]
- b) Random values between -0.05 and 0.05[3]
- c) Random values between 0 and 1 multiplied by a value of 0.01, but I changed that to 0.1[1]
- d) Random values between 0 and 1 divided by the square root of the number of inputs, but I changed the range of values to be between [-0.5, 0.5] or [-0.05, 0.05][1]

#### - Learning Rate:

The initial learning rate was chosen randomly, starting by a value of 1 and by using trial and error reducing that to 0.1, 0.01 and 0.001. Bigger values, i.e, 1, tended to give errors when calculating the cost, where it was calculated as nan, so I excluded 1 as a learning rate after the first experiment.

#### - Learning Rate Decay:

There were several methods suggested by research for choosing a value of learning rate decay[1], I chose two of these methods:

- a) Fixed decay by a value of 0.5 every 5 or 100 iterations

b) Constant decay by a value of 0.5 if cost does not improve after 50 or 100 iterations.

**- Epochs:**

The number of iterations the algorithm has to go through to optimize the weights and bias was chosen at random. I switched between 100 and 1000 epochs. I tried larger values of 10000, 20000 epochs but the improvement of model performance was negligible and did not justify the compute power used to run the iterations.

Experiment Number	Initial weights	Initial learning Rate	LR decay*	Number of Epochs	Cost at last epoch	Accuracy
01	$W = [-0.5, 0.5]$ [4]	1	lr /=2, epochs %5	100	nan	64%
02	$W = [-0.5, 0.5]$ [4]	0.1	lr /=2, epochs %5	100	2.04	56%
03	$W = [-0.5, 0.5]$ [4]	0.01	lr /=2, epochs %5	1000	2.80	40%
04	$W = [-0.05, 0.05]$ [3]	0.1	lr /=2, epochs %5	100	0.58	64%
05	$W = [-0.05, 0.05]$ [3]	0.1	lr /=2, epochs %5	1000	0.58	64%
06	$W = [-0.05, 0.05]/\sqrt{n}$ [3][1]	0.1	lr /=2, epochs %5	100	0.56	58%
07	$W = [-0.5, 0.5]/\sqrt{n}$ [4][1]	0.1	lr /=2, epochs %5	100	0.50	60%
08	$W = [-0.05, 0.05]/\sqrt{n}$ [3][1]	0.01	lr /=2, epochs %50**	1000	0.41	64%
09	$W = [0, 1]/\sqrt{n}$ [1]	0.1	lr /=2, epochs %100***	1000	nan	70%
10	$W = [0, 1]/\sqrt{n}$ [1]	0.01	lr /=2, epochs %100***	1000	0.12	72%
11	$W = 0.1 * [0, 1]$ [1]	0.01	lr /=2, epochs %50****	1000	0.14	76%
12	$W = [-0.5, 0.5]/\sqrt{n}$ [4][1]	0.1	lr /=2, epochs %100	1000	0.14	78%

\* Learning rate decays every few iterations[1]

\*\* If cost does not improve by 0.5 for 50 iterations

\*\*\* If cost does not improve by 0.001 for 100 iterations

\*\*\*\* If cost does not improve by 0.025 for 50 iterations

## 4 Results of the algorithm on the test dataset

Using the modified weights and bias from the previous training process to predict the output of a new "test" dataset achieved satisfactory results. I also collected a few pictures that were not in the dataset, then I preprocessed them as done previously with the train and test sets, i.e., I flattened then normalized them. I used them for testing the modified weights.

Giving the model more iterations to learn and train on images improved the overall performance as well as adjusting the learning rate to be adaptive, which allowed the model to take smaller steps when nearing the optimal/minimum cost, which aided in increasing the algorithm's accuracy and not overshooting and missing the optima.


Below are some screenshots of the output of testing on new images.

```
# Example of a picture
index = int(np.random.randint(low=0, high=49))
plt.imshow(testSetX[index])
plt.show()

class_name = [0, 0]

if round(test_pred.ravel()[index]) == 0:
    class_name[0] = 'non-cat'
elif round(test_pred.ravel()[index]) == 1:
    class_name[1] = 'cat'

print ("y = " + str(testSetY[:, index]) +
      ", predicted class= " + str(round(test_pred.ravel()[index])) +
      ", it's a " + class_name[int(round(test_pred.ravel()[index]))] + " picture.")
```



```
y = [1], predicted class= 1.0, it's a 'cat' picture.
```




```
# Example of a picture
index = int(np.random.randint(low=0, high=49))
plt.imshow(testSetX[index])
plt.show()

class_name = [0, 0]

if round(test_pred.ravel()[index]) == 0:
    class_name[0] = 'non-cat'
elif round(test_pred.ravel()[index]) == 1:
    class_name[1] = 'cat'

print("y = " + str(testSetY[:, index]) +
      ", predicted class= " + str(round(test_pred.ravel()[index])) +
      ", it's a '" + class_name[int(round(test_pred.ravel()[index]))] + "' picture.")
```




y = [1], predicted class= 1.0, it's a 'cat' picture.

```
# Example of a picture
index = int(np.random.randint(low=0, high=49))
plt.imshow(testSetX[index])
plt.show()

class_name = [0, 0]

if round(test_pred.ravel()[index]) == 0:
    class_name[0] = 'non-cat'
elif round(test_pred.ravel()[index]) == 1:
    class_name[1] = 'cat'

print("y = " + str(testSetY[:, index]) +
      ", predicted class= " + str(round(test_pred.ravel()[index])) +
      ", it's a '" + class_name[int(round(test_pred.ravel()[index]))] + "' picture.")
```




y = [0], predicted class= 0.0, it's a 'non-cat' picture.

```
# Example of a picture
index = int(np.random.randint(low=0, high=49))
plt.imshow(testSetX[index])
plt.show()

class_name = [0, 0]

if round(test_pred.ravel()[index]) == 0:
    class_name[0] = 'non-cat'
elif round(test_pred.ravel()[index]) == 1:
    class_name[1] = 'cat'

print ("y = " + str(testSetY[:, index]) +
        ", predicted class= " + str(round(test_pred.ravel()[index])) +
        ", it's a '" + class_name[int(round(test_pred.ravel()[index]))] + "' picture.")
```



y = [1], predicted class= 1.0, it's a 'cat' picture.

```
image_1 = plt.imread(imgs[0])
plt.imshow(image_1)

# Flatten and normalize image
image_1FN = image_1.reshape(-1).T / 255
print(image_1FN.shape)

# Use optimized weights for predicting the output
test_1 = sigmoid(np.dot(W_mod.T, image_1FN) + b_mod)

class_test_1 = [0, 0]

if round(test_1[0]) == 0:
    class_test_1[0] = 'non-cat'
elif round(test_1[0]) == 1:
    class_test_1[1] = 'cat'

print ("y = [0]" +
        ", predicted class= " + str(round(test_1[0])) +
        ", it's a '" + class_test_1[int(round(test_1[0]))] + "' picture.")
```

(12288,)  
y = [0], predicted class= 0.0, it's a 'non-cat' picture.



```
image_2 = plt.imread(imgs[1])
plt.imshow(image_2)

# Flatten and normalize image
image_2FN = image_2.reshape(-1).T / 255
print(image_2FN.shape)


# Use optimized weights for predicting the output
test_2 = sigmoid(np.dot(W_mod.T, image_2FN) + b_mod)

class_test_2 = [0, 0]

if round(test_2[0]) == 0:
    class_test_2[0] = 'non-cat'
elif round(test_2[0]) == 1:
    class_test_2[1] = 'cat'

print ("y = [0]" +
      ", predicted class= " + str(round(test_2[0])) +
      ", it's a '" + class_test_2[int(round(test_2[0]))] + "' picture.")
```

```
(12288,)
y = [0], predicted class= 1.0, it's a 'cat' picture.
```



```

image_3 = plt.imread(imgs[2])
plt.imshow(image_3)

# Flatten and normalize image
image_3FN = image_3.reshape(-1).T / 255
print(image_3FN.shape)

# Use optimized weights for predicting the output
test_3 = sigmoid(np.dot(W_mod.T, image_3FN) + b_mod)

class_test_3 = [0, 0]

if round(test_3[0]) == 0:
    class_test_3[0] = 'non-cat'
elif round(test_3[0]) == 1:
    class_test_3[1] = 'cat'


print ("y = [0]" +
      ", predicted class= " + str(round(test_3[0])) +
      ", it's a " + class_test_3[int(round(test_3[0]))] + " picture.")

```

```

(12288,)
y = [0], predicted class= 0.0, it's a 'non-cat' picture.

```



## 5 Marking Scheme

COURSEWORK 1a			
%1	QUESTIONS	S_MARK	FINAL
1	Did you code the entire algorithm?	x	
2	Did you vectorise your code?	x	
3	Is your code runnable?	x	
4	Did you collect data to measure convergence?	x	
5	Did you create a report explaining what you did?	x	
Total =			

## 6 Conclusions

Using the biological concept of a neuron and training it using backpropagation allowed for solving complex problems that were not easy to solve using rule-based systems, nor simple machine learning statistical functions. Artificial neural networks perform well in problems that relate to computer vision where a huge input is expected in the form of pixel data. I discussed in this report the building block of ANNs which is a single neuron and how to train it to classify cat images.

## 7 Appendices

### 7.1 Program Code

```
1
2 # coding: utf-8
3
4 # # Gradient Descent
5 #
6 # **F21BC Coursework 1**
7 #
8 # <sub>Name: **Heba El-Shimy**</sub>
9 # <br>
10 # <sub>Based on code obtained from coursework specification report ,
    written by **Dr. Marta Vallejo**</sub>
11
12 # In [1]:
13
14
15 # Import libraries
16 import numpy as np
17 import matplotlib.pyplot as plt
18 import h5py
19 import os
20 from math import sqrt
21
22
23 # In [2]:
24
25
26 np.set_printoptions(precision=2)
27
28
29 # In [3]:
30
31
32 # Loading the dataset
33 os.getcwd()
34 train_dataset = h5py.File('trainCats.h5', "r")
35 trainSetX = np.array(train_dataset["train-set-x"][:]) # your train set
    features
36 trainSetY = np.array(train_dataset["train-set-y"][:]) # your train set
    labels
37 trainSetY = trainSetY.reshape((1, trainSetY.shape[0]))
38
39 test_dataset = h5py.File('testCats.h5', "r")
40 testSetX = np.array(test_dataset["test-set-x"][:]) # your test set
    features
```

```

41 testSetY = np.array(test_dataset["test_set_y"][:]) # your test set
    labels
42 testSetY = testSetY.reshape((1, testSetY.shape[0]))
43
44 classes = np.array(test_dataset["list_classes"][:]) # the list of
    classes
45
46
47 # In [4]:
48
49
50 # Example of a picture
51 index = 20
52 plt.imshow(trainSetX[index])
53 plt.show()
54 print("y = " + str(trainSetY[:, index]) + ", it's a '" + classes[np.
    squeeze(trainSetY[:, index]).decode("utf-8") + "' picture.")
55
56
57 # In [5]:
58
59
60 # Images dimensions
61 print(trainSetX.shape)
62 print(testSetX.shape)
63
64 print('Image dimensions: {}px x {}px '.format(trainSetX.shape[1],
    trainSetX.shape[2]))
65 print('Image channels: {}'.format(trainSetX.shape[-1]))
66 print('Number of training examples: {} images'.format(trainSetX.shape
    [0]))
67 print('Number of test examples: {} images'.format(testSetX.shape[0]))
68
69
70 # In [6]:
71
72
73 # Flatten the pictures
74 # Applying (num_pixel x num_pixel x num_channels)
75 trainSetXF= trainSetX.reshape(trainSetX.shape[0], -1).T
76 testSetXF = testSetX.reshape(testSetX.shape[0], -1).T
77
78 print('Shape of training data after flattening: {}'.format(trainSetXF.
    shape))
79 print('Shape of test data after flattening: {}'.format(testSetXF.shape)
    )
80
81
82 # In [7]:

```

```

83
84
85 # Normalize images
86 # Applying (pixel_value/255)
87 trainSetXFN = trainSetXF / 255
88 testSetXFN = testSetXF / 255
89
90 print('Shape of training data after normalizing: {}'.format(trainSetXFN
    .shape))
91 print('Shape of test data after normalizing: {}'.format(testSetXFN.
    shape))
92
93 print('First row of training data before normalizing: \n{}\n'.format(
    trainSetXF[0]))
94 print('First row of training data after normalizing: \n{}\n'.format(
    trainSetXFN[0]))
95
96 print('First row of test data before normalizing: \n{}\n'.format(
    testSetXF[0]))
97 print('First row of test data after normalizing: \n{}\n'.format(
    testSetXFN[0]))
98
99
100 # In[8]:
101
102
103 # Network Topology
104 print('Number of input units: {}'.format(trainSetXFN.shape[0]))
105 print('Number of outputs: {}'.format(classes.shape[0]))
106
107
108 # In[336]:
109
110
111 # Initialize weights
112 W = np.random.uniform(low=-0.5, high=0.5, size=(trainSetXFN.shape[0],
    1)) / sqrt(trainSetXFN.shape[1])
113
114 print('Shape of weights matrix: {}'.format(W.shape))
115 print('Range of values in weights matrix = [{} - {}]'.format(W.min(), W
    .max()))
116 print('First (only) column in weights matrix: \n{}'.format(W))
117
118
119 # In[26]:
120
121
122 # Initialize biases
123 b = np.zeros([1, ])

```



```
124
125 #print('Shape of bias vector: {}'.format(b.shape))
126 print('First value in bias vector: \n{}'.format(b))
127
128
129 # In [21]:
130
131
132 # Activation function
133 # Sigmoid
134
135 def sigmoid(z):
136     """
137     Compute sigmoid function
138     @param z: value to compute sigmoid for (WX + b)
139     """
140
141     a = np.zeros([1, 1])
142     a = 1 / (1 + np.exp(-z))
143
144     return a
145
146
147 # In [22]:
148
149
150 # Cost calculation
151 # Cross-Entropy as the loss function
152
153 def cost(a, y):
154     """
155     Compute loss function
156     @param a: predicted label
157     @param y: actual label
158     """
159
160     L = np.sum((y * np.log(a)) + ((1 - y) * np.log(1 - a)))
161
162     J = (-1 / y.shape[1]) * L
163
164     return J
165
166
167 # In [338]:
168
169
170 # Training the neuron
171
172 W_mod = np.copy(W) # modified weights matrix
```

```

173 b_mod = np.copy(b) # modidied bias vector
174 lr = 0.1 # learning rate
175 epochs = 1000 # number of iterations
176 costs = [] # store all calculated costs
177
178 # Training iterations
179 for i in range(epochs):
180     J = 0
181     dW = np.zeros(W_mod.shape)
182     db = b_mod
183
184     z = np.dot(W_mod.T, trainSetXFN) + b_mod
185     predicted_labels = sigmoid(z)
186
187     J = cost(predicted_labels, trainSetY)
188     dW = (1 / trainSetY.shape[1]) * np.dot(trainSetXFN, (
189         predicted_labels - trainSetY).T)
190     db = (1 / trainSetY.shape[1]) * np.sum((predicted_labels -
191         trainSetY), axis=1)
192
193     costs.append(J)
194     # learning rate decay
195     if i > 0:
196         if i % 100 == 0: #and abs(costs[i-50] - J) <= 0.5:
197             lr /= 2
198
199     W_mod = W_mod - (lr * dW)
200     b_mod = b_mod - (lr * db)
201     print('Learning rate: {}'.format(lr))
202     print('Iteration {} \t ==> \t Cost: {:.2f} \n'.format(i, J))
203
204 # In [339]:
205
206 # Test the model on the test set
207 test_pred = sigmoid(np.dot(W_mod.T, testSetXFN) + b_mod)
208
209 # Measure model accuracy
210 count = 0
211 for pred, label in zip(test_pred.ravel(), testSetY.ravel()):
212     if round(pred) == label:
213         count += 1
214
215 accuracy = float(count) / testSetY.shape[1] * 100
216 print('Model\'s accuracy = {:.2f}%'.format(accuracy))
217
218
219 # In [343]:

```

```
220
221
222 # Example of a picture
223 index = int(np.random.randint(low=0, high=49))
224 plt.imshow(testSetX[index])
225 plt.show()
226
227 class_name = [0, 0]
228
229 if round(test_pred.ravel()[index]) == 0:
230     class_name[0] = 'non-cat'
231 elif round(test_pred.ravel()[index]) == 1:
232     class_name[1] = 'cat'
233
234 print ("y = " + str(testSetY[:, index]) +
235        ", predicted class = " + str(round(test_pred.ravel()[index])) +
236        ", it's a '" + class_name[int(round(test_pred.ravel()[index]))]
237        + "' picture.")
238
239 # In[405]:
240
241
242 import glob
243 import cv2
244 import os
245 import matplotlib.image as mpimg
246
247 # Experimenting on collected images
248
249 img_dir = "./images"
250 img_files = glob.glob(os.path.join(img_dir, '*.jpg'))
251 imgs = []
252
253 for img in img_files:
254     imgs.append(img)
255
256
257 # In[406]:
258
259
260 image_1 = plt.imread(imgs[0])
261 plt.imshow(image_1)
262
263 # Flatten and normalize image
264 image_1FN = image_1.reshape(-1).T / 255
265 print(image_1FN.shape)
266
267 # Use optimized weights for predicting the output
```

```

268 test_1 = sigmoid(np.dot(W_mod.T, image_1FN) + b_mod)
269
270 class_test_1 = [0, 0]
271
272 if round(test_1[0]) == 0:
273     class_test_1[0] = 'non-cat'
274 elif round(test_1[0]) == 1:
275     class_test_1[1] = 'cat'
276
277 print ("y = [0]" +
278        ", predicted class=" + str(round(test_1[0])) +
279        ", it's a '" + class_test_1[int(round(test_1[0]))] + "' picture
280        ".")
281
282 # In[407]:
283
284
285 image_2 = plt.imread(imgs[1])
286 plt.imshow(image_2)
287
288 # Flatten and normalize image
289 image_2FN = image_2.reshape(-1).T / 255
290 print(image_2FN.shape)
291
292 # Use optimized weights for predicting the ouput
293 test_2 = sigmoid(np.dot(W_mod.T, image_2FN) + b_mod)
294
295 class_test_2 = [0, 0]
296
297 if round(test_2[0]) == 0:
298     class_test_2[0] = 'non-cat'
299 elif round(test_2[0]) == 1:
300     class_test_2[1] = 'cat'
301
302 print ("y = [0]" +
303        ", predicted class=" + str(round(test_2[0])) +
304        ", it's a '" + class_test_2[int(round(test_2[0]))] + "' picture
305        ".")
306
307 # In[408]:
308
309
310 image_3 = plt.imread(imgs[2])
311 plt.imshow(image_3)
312
313 # Flatten and normalize image
314 image_3FN = image_3.reshape(-1).T / 255

```

```
315 print(image_3FN.shape)
316
317 # Use optimized weights for predicting the output
318 test_3 = sigmoid(np.dot(W_mod.T, image_3FN) + b_mod)
319
320 class_test_3 = [0, 0]
321
322 if round(test_3[0]) == 0:
323     class_test_3[0] = 'non-cat'
324 elif round(test_3[0]) == 1:
325     class_test_3[1] = 'cat'
326
327 print ("y = [0]" +
328        ", predicted class=" + str(round(test_3[0])) +
329        ", it's a '" + class_test_3[int(round(test_3[0]))] + "' picture
        .")
```

## 7.2 How to run the code

The code was written in a Jupyter Notebook as it is the best medium for viewing the results of fragments of code and iterate and test. To view the notebook, there should be installed on your computer a scientific python package, for example [Anaconda](#).

Download the notebook accompanied by this report and in a terminal window (unix) or cmd (windows), navigate to the folder that contains the recently downloaded `.ipynb` file and type `jupyter notebook`. This will spin up a local server on port 8888. You can view the notebook from your browser by navigating to <http://localhost:8888> and clicking on the notebook name `GradientDescent.ipynb`.

Also, a copy of the notebook with the outputs has been saved as a `.html` file with the name `GradientDescent.html` and has been uploaded with this report. It can be viewed it in any browser.

## References

- [1] Cs231n convolutional neural networks for visual recognition, 2018.
- [2] Adel El-Shahat. Introductory chapter: Artificial neural networks. In Adel El-Shahat, editor, *Advanced Applications for Artificial Neural Networks*, chapter 1. IntechOpen, Rijeka, 2018.
- [3] Mercedes Fernández-Redondo and Carlos Hernandez Espinosa. Weight initialization methods for multilayer feedforward, 01 2001.
- [4] Marta Vallejo. Biologically inspired computation, 2018.